

Persisting Intentions

ALFRED R. MELE

Florida State University

Al is nearly finished sweeping his kitchen floor when he notices, on a counter, a corkscrew that should be put in a drawer. He intends to put the corkscrew away as soon as he is finished with the floor; but by the time he returns the broom and dustpan to the closet, he has forgotten what he intended to do. Al knows (or has a true belief) that there is something he intended to do now in the kitchen. He gazes around the room and tries to recall what it was. Within a minute or so, without seeing the corkscrew, Al recalls. He puts the corkscrew away.

Did Al have an intention to put the corkscrew away that persisted from the time he acquired the intention until he put the corkscrew away? That is one of my guiding questions in this article. Perhaps surprisingly, an attempt to come to grips with it leads quickly to some largely uncharted territory in the philosophy of mind and action. Although considerable attention has been paid to nonoccurrent beliefs and desires, little attention has been paid to nonoccurrent intentions; and apparent differences between occurrent and nonoccurrent intentions, as I will explain, bear directly on my question about Al and the corkscrew.

To make things more manageable than they would otherwise be, I make three assumptions. First, intentions exist. Second, anyone who intends to *A* has an intention to *A*.¹ Third, some causal account or other of the production of intentional actions is correct, and, more specifically, some such account according to which whenever agents act intentionally, at least one of the following does important causal work: some intention of theirs; the acquisition or persistence of some intention of theirs; or the neural correlate of one or more of the preceding items. I have defended an account of this kind elsewhere, as have others.² To do so again here would leave no space

for exploring this article's topic. With the aim of contributing to our growing understanding of the place of intentions in the lives of intelligent agents, I develop accounts of what it is for an intention to be occurrent and of the persistence of an occurrent intention, and I do something similar for relatively short-term "standing" intentions.

1. Kinds of Intention

In what does the persistence of an intention consist? One way to approach this question is to develop an account of intentions and then ask about the persistence of intentions, so conceived. However, developing an account of intentions is a project in its own right, I have undertaken a major part of the project elsewhere, and my simply appealing at this point to the position I defended would understandably evoke some grumbling.³ So, for now, I encourage readers to understand the term "intention" in whatever way they normally understand it, provided that their understanding of it is consistent with the truth of the three assumptions I stated. (Readers whose understanding of intention is inconsistent with the truth of those assumptions have the option of setting this article aside right now, but that might be rash.)

A common distinction between *proximal* and *distal* intentions is compatible with various alternative ways of understanding intentions.⁴ In the simplest case, a proximal intention to *A* is an intention to *A* now: for example, an intention to answer one's phone now. Proximal intentions also include intentions to continue doing something one is presently doing and intentions to start *A*-ing (e.g., start running a mile) straightaway. Distal intentions are intentions agents have at a time for (and only for) action in the nonimmediate future: for example, the intention I have now, at noon, to take my daughter to a 7:00 movie tonight.

Some intentions have both proximal and distal aspects. For example, Al may have an intention to run a mile without stopping, beginning now. (He estimates that the deed will take six minutes.) I call such an intention a *mixed* intention. An intention of this kind specifies something to be done now and something to be done later. Just as there is no precise point of demarcation between men who count as bald and men who do not, there is no precise point of demarcation between intentions that count as proximal and intentions that count as mixed.

If intentions of all three kinds persist, do they persist in the same way? Attention to a familiar distinction between *occurrent* and *standing* desires or wants (Goldman 1970, 86–88; see Alston 1967, 402) will prove useful as background for an attempt to answer this question.⁵ Alvin Goldman writes: "An occurrent want is a mental event or mental process; it is a 'going on' or 'happening' in consciousness. A standing want. . . is a disposition. . . to have an occurrent want, a disposition that lasts with the agent for a reasonable length of time" (86). For example, we can "say that John wants to be president

of his company because he has a *disposition* to have occurrent wants to this effect. . . . This disposition is present. . . . though it is not being manifested” (Goldman 86). A distinction of this kind is motivated partly by evidence about the ordinary notion of desire or wanting provided by common attributions of desires or wants. Beth, a graduate student, is deeply committed to pursuing a career in biology. As the verb “want” is standardly used, “Beth wants a career in biology” is true even while she is thoroughly absorbed in a conversation about her father’s financial problems or dreamlessly sleeping. When Beth’s supervisor finds himself inclined to tell a prospective employer that she definitely wants a career in biology, he need not phone her to learn whether she is awake before he can be confident that he will be speaking the truth. The quoted sentence is also true when Beth is crafting a cover letter for her job applications. In the latter case, but not the former, Beth has an occurrent want for a career in biology.

Goldman’s way of drawing the distinction is problematic. An important part (at least) of what the distinction is supposed to mark is the difference between (1) desires to *A* (or for *X*) that are wholly constituted by dispositions to have desires to *A* (or for *X*) that manifest these dispositions and (2) desires that are functionally more closely connected to action than such dispositions are (Alston 1967, 402, Goldman 1970, 88). For example, the characteristic functional role of Beth’s standing desire for a career in biology is to contribute to her having occurrent desires for a career in biology. Her standing desire’s functional connection to action features its contributing to her having occurrent desires that manifest it. But not all desires with a more direct connection to action are present to consciousness—at least if common sense and clinicians may be trusted.⁶ Carl might have spoken to Debra as he did because he wanted to hurt her feelings, even though he was not conscious at the time of a desire to do that. Carl might not have had a standing desire to hurt Debra’s feelings. But if he did have one, and if it played the role characteristic of standing desires, it was manifested in a desire of another kind to hurt her feelings—an occurrent desire to do that. (If Carl did not have a standing desire to hurt Debra’s feelings, then, of course, any occurrent desire to hurt her feelings that he had did not manifest a standing desire to do that.) Goldman implausibly makes being “a ‘going on’ or ‘happening’ in consciousness” an essential feature of occurrent wants or desires. An agent’s being conscious of a desire is not required for the desire’s playing a role in action production that suffices for its not being a standing desire and therefore suffices for its being an occurrent desire, if the distinction between standing and occurrent desires is exhaustive. I should add that occurrent desires, as I understand them (Mele 2003, 30–33), are not occurrences or events (“in consciousness” or otherwise); they are states.

As I understand *standing* desires, intentions, and beliefs, they are wholly constituted by dispositions of certain kinds to have corresponding occurrent desires, intentions, and beliefs. For example, a standing belief that *p* is

wholly constituted by a disposition of a certain kind to have occurrent beliefs that *p*. These dispositions have no explicit representational content. They are accorded a content-supplying description in virtue of the content of the occurrent attitudes that manifest them. Having no explicit representational content, standing desires, intentions, and beliefs are not *explicit attitudes*.⁷ Rather, they are dispositions to have explicit attitudes of a certain kind. They are, one may say, “implicit” attitudes.⁸

If there are proximal intentions, if they tend to be fit for execution, and if they play an important role in the production of intentional actions, then, in typical instances, agents begin executing their proximal intentions almost as soon as they acquire them.⁹ If standing intentions to *A* are wholly constituted by dispositions to have occurrent intentions to *A*, then, because such dispositions are not executable, no executed intention is a standing intention. And simply in virtue of its being executed at a time, a proximal intention has a claim to being an occurrent intention at that time. (If the distinction between standing and occurrent intentions is exhaustive, then all executed intentions are occurrent intentions.)

What about distal intentions? Although many of them are intentions for prospective actions to be undertaken hours, days, or years later, their acquisition may have some significant practical effects within seconds or milliseconds. Recall my intention to take my daughter to a 7:00 movie tonight. What happened is that my daughter, Angela, called me at work shortly before noon to invite me to a 7:00 movie. At noon, I decided to do it. In so deciding, I formed an intention to do it. (Like many others, I take deciding to *A* to be a mental action of forming an intention to *A*.)¹⁰ And then, as I continued my telephone conversation, I wrote a note to remind myself to pick Angela up at 6:30 and I put it in my pocket. My intention to take Angela to a 7:00 movie seemingly played an important role at the time in generating my intention to write the note. If the “movie” intention is at work in motivating the “note” intention, it has a claim to being an occurrent intention at the time it is so functioning.

I remembered the movie plan without even having to be reminded by my note, and I executed it without a hitch. If my intention to take Angela to a 7:00 movie—to *M*, for short—persisted from noon into the evening, in what did its persistence consist? Did it persist as an *occurrent intention* throughout that span of time?

Return to Goldman’s account of an occurrent want as “a ‘going on’ or ‘happening’ in consciousness.” If occurrent intentions were to be defined as intentions of which the agent is conscious, it would be clear that my intention to *M* did not persist as an occurrent intention from noon into the evening. Surely, I was not conscious of that intention throughout the entire span of several hours. This account of occurrent intentions is implausible for the same reason that Goldman’s account of occurrent wants is implausible: Carl might have spoken to Debra as he did because he intended to hurt her feelings, even

though he was not conscious at the time of his intention to do that.¹¹ Even so, the verdict that it yields about any intention of mine to *M* that persisted from noon into the evening may be correct.

As I understand occurrent intentions, that agents' intentions to *A* (or their persistence, their acquisition, or neural correlates thereof) are nondeviantly at work at a time in initiating, sustaining, or guiding *A*-ings of theirs that are being initiated at that time or are in progress at the time is conceptually sufficient for the intentions' being *occurrent* intentions at that time, provided that the contribution the intentions to *A* are making is not made by way of their being manifested in other intentions to *A*.¹² (Henceforth in this article, I suppress parenthetical clauses of the sort that appear in the preceding sentence and write simply in terms of *intentions* being at work.) But this is not also a necessary condition for an intention's being an occurrent intention at a time. For example, one need not actually *A* at a time in order for one's intention to *A* to be occurrent at that time: one might try to *A* and fail. Furthermore, as I mentioned, an occurrent intention to *A* may be at work in producing intentions and actions that are preparatory to one's *A*-ing, as in the case of my writing a note to remind myself about the movie. Being aware of an intention to *A* that one has—or being aware that one intends to *A*—may also be regarded as a conceptually sufficient condition of the intention's being occurrent at the time, provided that what one is aware of is not something that is wholly constituted by a disposition to have occurrent intentions to *A* and provided that one's awareness that one intends to *A* is not merely awareness that one has a disposition—even a very strong one—to acquire occurrent intentions to *A*.¹³ (A reflective smoker may be aware that he is strongly disposed to have occurrent desires and occurrent intentions to smoke after meals and wonder why he has no desire and no intention to smoke now, after finishing a meal.) For example (with the proviso just identified in place), my being aware at 3:00 that I intend to take Angela to a 7:00 movie may be treated as sufficient for my having at the time an occurrent intention to take her to the movie.

If there are both occurrent and nonoccurrent intentions to *A*, then any intention to take Angela to a 7:00 movie that I have from noon into the evening is a more promising candidate for being an intention of the latter kind than for being an intention of the former kind. One can try to make a case for that intention's being at work in some relevant ways throughout the whole span of time, but the kind of case one can make for that can, in principle, be made about any putative nonoccurrent intention in normal circumstances. One who deems the case successful, should see the distinction between occurrent and nonoccurrent intentions as collapsing.

If standing intentions to *A* are understood on the model of a familiar conception of standing wants, they are wholly constituted by dispositions to have occurrent intentions to *A*. Now, not all dispositions to have occurrent intentions to *A* are promising candidates for being intentions. A hypnotist

implants in Erwin a disposition to acquire an occurrent intention to open a window whenever he hears the word “open” in the presence of a window. It seems that, as yet, Erwin has no new window-opening intention—not even a standing one. Rather he is disposed to have window-opening intentions in certain circumstances. Of course, this is consistent with its being true that some dispositions to have occurrent intentions to *A* are standing intentions to *A*. If after I decided to take Angela to a 7:00 movie, I had an intention to do that that persisted for several hours, perhaps it was a standing intention constituted by a certain disposition to have occurrent intentions to do that. I return to standing intentions in section 3.

I have discussed connections between occurrent and nonoccurrent intentions, on the one hand, and proximal and distal intentions on the other. What about mixed intentions: for example, Al’s intention to run a mile, beginning now? If Al’s acquiring that intention straightaway initiates his beginning to run a mile, the intention has a claim to being an occurrent intention at that time.

2. Persisting Occurrent Intentions

In section 1, I offered, in effect, two different sufficient conditions for its being the case that an intention to *A* is an occurrent intention at a time: (1) the intention is suitably at work at that time in producing relevant intentional actions or in producing items appropriate for the production of relevant intentional actions; (2) at that time the agent is aware that he intends to *A* or aware of his intention to *A* and this intention is not wholly constituted by a disposition to have occurrent intentions to *A*.¹⁴ I suggest now that a necessary condition for its being the case that an intention is an occurrent intention at a time is that either 1 or 2 is true of it.¹⁵ As shorthand for the awareness condition specified in 2, I use the expression “conscious intention.”¹⁶ Given this account of what it is for an intention to be occurrent, an occurrent intention to *A* persists as an occurrent intention to *A* as long as it is suitably at work or is a conscious intention that satisfies condition 2.

Now, understanding “occurrent” in the expression “occurrent intention” is one thing and understanding “intention” in that same expression is another. To be in a position to defend an interesting account of what it is for one and the same occurrent intention to be at work—or to be an object of an agent’s awareness—over a span of time, one would need to defend an account of what it is for an occurrent item to be an *intention*.

In the present section, I construct an account of persisting occurrent intentions. I have defended an account of the nature and functions of non-standing intentions elsewhere (Mele 1992), and I lack the space to do so again here. Given that having some substantive account of such intentions on the table would facilitate moving forward, I offer a brief summary of the account. What is at issue now is what it is for an occurrent item to be an intention.

Occurrent intentions, on my view, are executive attitudes toward plans (Mele 1992, chs. 8–11).¹⁷ Like occurrent beliefs and desires, occurrent intentions have representational content. The representational content of an occurrent intention is an action-plan. In the limiting case, the plan-component of an occurrent intention has a single “node.” It is, for example, a prospective representation of one’s taking a vacation in Lisbon next winter that includes nothing about means to that end nor about specific vacation activities. Often, intention-embedded plans are more complex. The intention to check her e-mail that Jan executed this evening incorporated a plan that included clicking on her e-mail icon, then typing her password in a certain box, then clicking on the “OK” box, and so on. An agent who successfully executes an occurrent intention is guided by the intention-embedded plan.¹⁸

Although the contents of intentions are plans, I follow the standard practice of using such expressions as “Jan’s intention to check her e-mail now” and “Ken intends to bowl tonight.” It should not be inferred from such expressions that the agent’s intention-embedded plan has a single node—for example, checking e-mail now or bowling tonight. Often, our expressions of an agent’s desires and intentions do not identify the full content of the attitude and are not meant to. Jan says, without intending to mislead, “Ken wants to bowl tonight,” knowing full well that what he wants is to bowl with her at MegaLanes tonight for twenty dollars a game until the place closes, as is their normal practice.

According to a popular view of occurrent representational attitudes—for example, Lara’s belief that *p*, Mel’s desire that *p*, Nora’s desire to *A*, Owen’s intention to *A*—one can distinguish between an attitude’s representational content and its psychological *orientation* (Searle 1983). Orientations include believing, desiring, and intending. On my view, the executive dimension of occurrent intentions is intrinsic to the attitudinal orientation *intending*. We can have a variety of attitudes toward plans: for example, we might admire plan *x*, be disgusted by plan *y*, and desire to execute plan *z*. To have the intending attitude toward a plan is to be settled (but not necessarily irrevocably) on executing it.¹⁹ The intending and desiring attitudes toward plans differ in that the former alone entails this settledness. Someone who desires to *A*, or to execute a certain plan for *A*-ing—even someone who desires this more strongly than he desires not to *A*, or not to execute that plan—may still be deliberating about whether to *A*, or about whether to execute the plan, in which case he is not settled on *A*-ing, or not settled on executing the plan.²⁰ Pat wants more strongly to respond in kind to a recent insult than to refrain from doing so, but, owing to moral qualms, she is deliberating about whether to do so. She is unsettled about whether to retaliate despite the relative strength of her desires (see Mele 1992, ch. 9).

On a standard view of desire, the psychological features of occurrent desires to *A* in virtue of which they contribute to intentional *A*-ings are their content and their strength. On my view of the contribution of occurrent

intentions to *A* to intentional *A*-ings, the settledness feature of intentions is crucial, and it is not capturable in terms of desire strength (and content), nor in terms of this plus belief (Mele 1992, 76–77 and ch. 9). Occurrent intentions to *A*, as I understand them, essentially encompass motivation to *A*, but without being reducible to a combination of desire and belief (Mele 1992, ch. 8). Part of what it is to be *settled* on *A*-ing is to have a motivation-encompassing attitude toward *A*-ing; lacking such an attitude, one lacks an element of a psychological commitment to *A*-ing that is intrinsic to being settled on *A*-ing, and therefore to having an occurrent intention to *A*.

There is a body of literature on belief constraints on intentions. Elsewhere, I have defended the view that *S* intends to *A* only if *S* lacks the beliefs that he will not *A* and that he probably will not *A* (Mele 1992, 146–51). If this is right (but, for the purposes of this article, I have no need to insist that it is), an agent may have an executive attitude toward a plan for *A*-ing without having an intention to *A*. For example, a golfer may have an executive attitude toward a plan for sinking a putt that he believes he probably will miss. But this is not to say that the attitude at issue is not an intention. It is plausibly deemed an intention to try to sink the putt.

Given the account of occurrent intentions just sketched, whether an occurrent intention to *A* persists depends, among other things, on whether the agent acquires a belief that he probably will not *A*. An agent who intends to unlock his office door may try, fail, and try again. After another attempt or two, he may come to believe that he probably will not unlock it and try again anyway. On the view sketched, his acquisition of this belief suffices for the nonpersistence of his intention to unlock the door. Now, even before he acquires the belief, his plan has probably changed. Originally, the plan was his normal plan for unlocking the door. Subsequently, it may have featured an uncharacteristic bit of jiggling. At that point, before he acquires his pessimistic belief, it is still true that he intends to unlock the door; but the change in plan might be deemed sufficient for the nonpersistence of his original intention to unlock the door. Earlier, he intended to unlock it in the normal way; now he intends to unlock it in another way.

Many persisting things change considerably over time. I have changed in various ways in the last several years, as has my house and all of the persisting things in it. In what ways can an occurrent intention change while it persists? If occurrent intentions are occurrent executive attitudes toward plans, an occurrent intention cannot persist while no longer being an executive attitude or while no longer having a plan as its content. Must an occurrent intention that one has at *t1* and an occurrent intention that one has at a later time *t2* encompass the *same plan* in order to be the same intention?

Recall that such expressions as “Jan’s intention to check her e-mail now” should not mislead us into thinking that the agent’s intention-embedded plan is simply to check e-mail now. I have a specific default route for driving home from work. Normally, when I start my car in the parking lot after work with

the intention of driving home, the plan-component of that intention includes driving along my default route, as I understand the contents of occurrent intentions. I also have several back-up routes. If I see that traffic is backed up along my normal route, I take an alternate route. When I start taking an alternate route, it is still true of me that I intend to drive home; but, I suggest, the intention to drive home that I have at this point is not the intention to drive home that I had earlier, owing to my having opted against a node in my original plan in favor of a course of action that was not in my plan.

When an agent is in the course of executing an intention, is any plan-change in the agent sufficient for one intention's being replaced by another? (Readers should not read "plan-change in an agent who is executing an intention" as "plan-change in an intention that is being executed." While I am executing an intention to drive home from work, I may change my plans about what project to start working on next week.) Drawing distinctions among some different kinds of plan-change that can occur while an agent is executing an intention will prove useful in attempting to answer this question. (I do not regard the kinds of plan-change to be considered here as exhaustive.) In each case, it is to be understood that the agent is in the process of executing an intention.

My opting not to do something I had planned to do (e.g., drive home along my default route) and to do something else instead (e.g., drive home along alternate route 1) is sufficient for plan-change. Call this *plan-change by substitution*.

Sometimes intentions that an agent is in the process of executing are very loose. Imagine that I have scheduled a meeting with a student, Zeke, for the purpose of talking him out of writing his dissertation on topic X and of guiding him toward topics that are more suitable for him. I know Zeke and his work well, and, because I have done this sort of thing many times and have been happy with the results, I count on myself to find a way of succeeding during our conversation without any advance planning. I start the meeting with an intention the plan-component of which may be fairly accurately described as follows: to persuade Zeke not to write on X; to identify, partly in light of our conversation, some good dissertation topics for him; and to persuade him to think seriously about which of those topics he would be most interested in pursuing. During the course of our conversation, comments on specific points strike me as called for; and I make them—intentionally, of course. As Zeke is explaining why he wants to write on X after having expressed some worries about the impact a protracted dissertation process will have on his financial situation, I decide to explain to him later why he is not likely to finish a dissertation on that topic in a reasonable amount of time and then to launch into an explanation of why he would have a much better chance of completing a good dissertation on various other topics in a timely fashion. After Zeke has finished his pitch for X, I do just that. This is an instance of *plan-change by addition*.

Sometimes, when I depart from my default route for driving home from work, I have a plan that includes turning right off of my default route at A Street and then either turning left at First Street or continuing down A Street, depending on which street is less congested. When I get to the First Street intersection, I make a decision or acquire an intention about whether to continue down A Street or to turn onto First Street. My becoming settled at this point about which to do is sufficient for plan-change. This is an instance of *plan-change by fork resolution*.

Regarding a particular case of plan-change by substitution, I have already suggested that the intention that was being executed does not survive the change, owing to the agent's having opted against a node in his original plan in favor of a course of action that was not in his plan. I suggest now that this is a general truth about plan-change of this kind.

What about plan-change by addition? I suggest that, in the illustration I sketched, my original intention survives, and I acquire what may be termed *implementing intentions* that coexist with the original one. I also suggest that the same is true of plan-change by fork resolution. In the example I offered, when I get to the First Street intersection, the following portion of my plan is most directly relevant: turn left at First Street if it is less congested than A Street; otherwise, continue driving on A Street. I do not opt against either of these conditional elements of my plan (that is, either conditional sub-plan): each conditional continues to be fine with me. (Analogy: one who believes a statement of the form "if p , then q " should not be led to reject it by the discovery that p is false.) Rather, I form or acquire an implementing intention—to turn left onto First Street, for example.²¹

Suppose that my suggestions about the preceding three kinds of plan-change are true. Then a plausible diagnosis of why it is that the original intention does not persist in the first kind of plan-change but does persist in the second and third kinds is that the first kind, unlike the other two, includes the agent's opting against some element of the plan-component of the original intention in favor of a course of action that is not included in that plan.

I asked whether an occurrent intention $N1$ that one has at $t1$ and an occurrent intention $N2$ that one has at a later time $t2$ need to have the same plan-component in order to be the same occurrent intention. If the correct answer is *yes*, we have a simple necessary condition for the persistence of an occurrent intention to A as an occurrent intention to A : (*OIP*) Occurrent intention $N1$ at $t1$ and occurrent intention $N2$ at $t2$ are the same occurrent intention only if they have the same plan-component. My suggestions about intention-change and intention-persistence in the face of plan-change are consistent with the truth of *OIP*. In plan-change by substitution, my earlier and later intentions "to A " have different plan-components, and my suggestion was that they are different intentions. In plan-change by addition and plan-change by fork resolution, the implementing intentions have different

plan-components than the original intention and are different intentions from it; and my suggestion that the original intention survives is consistent with the truth of the claim that it does so only if its plan-component has not changed.

Are there grounds for rejecting *OIP*? My default route home from work is about two miles long. On a normal day, when I am about a third of the way home from work, with what intention am I driving? Someone may claim that it is an intention whose plan-component starts at my current location; and, of course, the plan-component of the intention with which I started the trip did not start there. This claim is a problem for *OIP* only in conjunction with the further claim that, despite the difference in plan-components, the later intention is the same intention as my initial one. This further claim should be rejected: surely, in an agent who knows that point A and point B are not the same point, an intention to drive from B to C is not the same intention as an intention to drive from A to C. But what about the initial claim? Do I acquire new driving-home intentions from moment to moment as I progress along my default route, or do I continue to execute an intention whose plan-component begins in the parking lot at work? I opt for the latter—more economical—hypothesis.

It may be suggested that cases in which there are only minor revisions in an agent's plan for *A*-ing as the agent is in the process of *A*-ing falsify *OIP*. Now, taking alternate route 1 home from work involves only a minor departure from my default route: it brings me to the main stretch of my default route only one block east of my normal entry to it, and I drive east on the main stretch anyway. Even so, by driving home along alternate route 1, I cannot execute an intention to drive home along my default route. So no intention I execute by driving home along alternate route 1 is the intention with which I started the trip in those cases in which I started the trip with an intention to drive home along my default route.

It may be replied that even if the plan-component of the intention with which I start my trip home is a representation of driving home along my default route, that intention does not depend for its identity on its plan-component and, accordingly, I may execute the intention while departing from the plan. This reply is more radical than it may sound. My position on intentions is guided by a common view of attitude individuation: if attitude A and attitude B do not have the same content, they are not the same attitude. In my view, as I mentioned, the contents of intentions are what I have been calling their plan-components. The radical reply needs to be backed up by a defense of an alternative way of understanding the contents of intentions or by an argument against the common view of attitude individuation.

Even if *OIP* is plausible and true, a biconditional version of it is implausible. Consider the following: (*OIPa*) Occurrent intention *N1* at *t1* and occurrent intention *N2* at *t2* are the same occurrent intention if and only if they have the same plan-component. Compare the occurrent intention I

had at noon to take Angela to a 7:00 movie (*occurrent intention 1*) with the occurrent intention I had to do that at 6:00 when I was thinking about when to leave my office to pick her up (*occurrent intention 2*). They seem to have the same plan-component. But, as I have explained, it is implausible that my occurrent intention to take Angela to a 7:00 movie persists as an occurrent intention from noon into the evening. I came to have occurrent intention 1 at noon. If, as it seems, occurrent intention 2 emerged much closer to 6:00, intentions 1 and 2 are not the same occurrent intention (even if they are tokens of the same occurrent intention-type). Also, given what I have said about belief constraints on intentions, I am committed to the possibility that two intentions may have the same plan component even though one is an intention to *A* and the other is instead an intention to try to *A*.

The following biconditional avoids both problems with *OIPa*:

OIPb. In an agent *S*, occurrent intention *N1* at *t1* and occurrent intention *N2* at *t2* are the same occurrent intention if and only if (1) they have the same plan-component, (2) there is no time between *t1* and *t2* at which no intention with that plan-component is suitably at work in *S* in producing intentional actions performed by *S* or in producing items appropriate for the production of intentional actions by *S* or is a conscious non-standing intention, and (3) between *t1* and *t2* *S* undergoes no change of belief that would suffice for an intention with that plan-component being a different intention from *N1*.²²

A version of *OIPb* asserting that the biconditional is a necessary truth is false if, for example, God can replace one intention with another that has the same plan-component quickly enough to permit the right-hand side of the biconditional to be satisfied even though a switch has taken place. I have no wish to argue about whether God can do this. Adding a condition to the effect that *S* has undergone no instantaneous and no nearly instantaneous intention-switching between *t1* and *t2*—the switched intentions having the same plan-component—would block the potential counterexample.²³

3. Persisting Short-Term Standing Intentions

The standing intentions that concern me in this article, as I mentioned, are relatively short-term dispositions. If there are standing intentions, an agent may perhaps have, for years, a standing intention to exercise every day or to finish writing a book on procrastination. My concern here is with shorter-term scenarios, such as that of Al and the 7:00 movie and Al and the corkscrew, if Al has standing intentions regarding the movie and the corkscrew in these scenarios. I return to my question about the latter scenario: Did Al have an intention to put the corkscrew away that persisted from the time he acquired the intention until he put the corkscrew away? In light of the preceding two sections, two additional questions should be asked. Did Al have a persisting occurrent intention to put the corkscrew away? Did he have a persisting

standing intention to do this? Recall that the span of time at issue is no more than a few minutes. Al temporarily forgot what it was that he intended to do next in the kitchen; but when he remembered, he put the corkscrew away.

On the account I have offered, an intention is an *occurrent* intention at a time if and only if (1) the intention is suitably at work at that time in producing relevant intentional actions or in producing items appropriate for the production of relevant intentional actions or (2) the intention is a conscious non-standing intention at that time. After putting the broom and dustpan away, Al paused because he forgot what he intended to do next. If Al had had an intention to put the corkscrew away then that was suitably at work at the time, he would have put it away then. Similarly, if he had had a conscious non-standing intention to put the corkscrew away then, he would have put it away then. Instead, he paused to try to recall what he intended to do next.

This leaves it open that Al has a persisting standing intention to put the corkscrew away, that he has a persisting intention to do this that is *occurrent* at some times but not *occurrent* at others, and that he has no persisting intention at all to do this. Which is it? Even while Al has temporarily forgotten what he intended to do, he may have a disposition to recall that he intended to put the corkscrew away, and that disposition may be a major part of a disposition Al has to have *occurrent* intentions to put it away. Now, the claim that something or someone has a disposition to *x* does not entail that it, he, or she, is more likely to *x* than not to *x*. When this is borne in mind, it should be clear that Al does have a disposition to recall what he intended to do next. If, in the circumstances, recalling this is sufficient to generate a proximal intention to put the corkscrew away, is Al plausibly regarded as having, for some time, a standing intention to put it away when he finishes sweeping the floor—a standing intention to *C*, for short?

Brief attention to *occurrent* and standing beliefs will shed some light on the issue. Robert Audi (1994) argues for a distinction between standing beliefs that *p* and mere dispositions to acquire *occurrent* beliefs that *p*. Most of the time, Fred has only a standing belief that Lubbock is in Texas; sometimes he has an *occurrent* belief that this is so. The standing belief is a disposition to have *occurrent* beliefs that Lubbock is in Texas. But is it plausible that all dispositions to have *occurrent* beliefs that *p* are beliefs—even standing ones—that *p*? Before you read this sentence, you had a disposition to believe *occurrently* that a man who has sixty-seven plastic pink flamingoes on his lawn is displaying more lawn ornaments, other things being equal, than a woman who has twenty-three plastic red bears on her lawn. But that disposition is very different from Fred's standing belief that Lubbock is in Texas. That Lubbock is in Texas is something that Fred learned, that he *occurrently* believed at times in the past, and that he remembers. Your relation to the fact about lawn ornaments is of a very different kind. Presumably, you encountered that fact here for the first time. Perhaps a necessary condition of its being true that a disposition to have *occurrent* beliefs that *p* is a standing

belief that *p* is that the person's acquiring an occurrent belief that *p* played a role in producing the disposition.

Is this supposed necessary condition also a sufficient condition? Is the following thesis true?

SB. A disposition to have occurrent beliefs that *p* is a standing belief that *p* if and only if the person's acquiring an occurrent belief that *p* played a role in producing the disposition.

Suppose that Fred's brother Ed, who also learned many years ago that Lubbock is in Texas, has, for the past ten years, consistently failed to recall where Lubbock is when the question arose, as it has several times. In each case, Ed's circumstances were normal: he was neither tired nor drunk, for example. However, he has a weak disposition to have occurrent beliefs that Lubbock is in Texas: there is a remote but genuine chance that persistent prompting by Fred, who spent a lot of time with Ed learning the locations of American towns when they were children, would result in Ed's remembering where Lubbock is, if Fred's prompting were to include his playing recordings of some cowboy songs that the boys listened to when trying to learn the names of Texas towns. Ed's disposition to have occurrent beliefs that Lubbock is in Texas satisfies the right-hand side of *SB*. But given how difficult it is for Ed to recall that Lubbock is in Texas, the judgment that it is false that he believes that Lubbock is in Texas and therefore false that he has a standing belief that Lubbock is in Texas is plausible.

Al's disposition to have occurrent intentions to *C* was produced in part by his acquiring an occurrent intention to *C*. So if there is a necessary condition for standing intentions that parallels the supposed necessary condition for standing beliefs that I mentioned, Al's disposition satisfies it. Now, the reason I gave for denying that Ed has a standing belief that Lubbock is in Texas can be recast as the claim that his disposition to have occurrent beliefs that this is so is *too weak*. Given the details of Al's case, was his disposition to have occurrent intentions to *C* too weak for him to have a standing intention to *C* during the period in question?

What we should think about this is linked to what we should think about an empirical issue. Suppose that whenever normally functioning human beings consciously acquire a distal intention to *A* that they do not proceed straight-away to revoke, there is an interesting internal effect. One possible effect is that the distal intention—an executive attitude toward a plan—is stored as an intention in an “intention buffer.” Another possible effect (among others) is an internal representation of the fact that the agent acquired a distal intention to *A*.

Suppose that the former effect is an actual effect. While the agent is aware that he intends to *A*, the stored distal intention is an occurrent intention. Suppose that even if the agent forgets that he intended to *A*, the stored intention

persists for a time. Because the intention is not itself a disposition to have occurrent intentions, it is not a standing intention. And if, during some span of time, it is neither at work in relevant ways nor a conscious intention, it is not an occurrent intention during that time. (Here, I am assuming the correctness of my account of what it is for an intention to be occurrent.) It is, then, an intention of a third kind—an intention that can exist while being neither a standing intention nor an occurrent intention. Call it a *3-intention*. If there are 3-intentions, and if, partly because he has a 3-intention to *A*, an agent has a disposition to acquire occurrent intentions to *A*, there are interesting grounds for the claim that that disposition is a standing intention to *A* even when the agent has forgotten that he intended to *A*. The grounding of the disposition partly in a persisting intention to *A*—a 3-intention—provides some reason to judge that the disposition counts as a standing intention to *A*.

Obviously, if the stored-intention effect that I have described is not an actual effect, this reason for counting the disposition as a standing intention dissolves; and the supposition that an internal representation of the fact that the agent acquired a distal intention to *A* is an actual effect does not bring this reason back. Now, I have suggested that Al may have a disposition to recall that he intended to *C* and that that disposition may be a major part of a disposition Al has to have occurrent intentions to *C*. If Al's disposition to recall that he intended to *C* is as weak as Ed's disposition to recall that Lubbock is in Texas and Al lacks a 3-intention to *C*, the claim that Al has a standing intention to *C* seems no more plausible than the claim that Ed has a standing belief that Lubbock is in Texas. Of course, Al may temporarily forget that he intended to *C* and still have a fairly strong disposition to recall this. Other things being equal, provided that, in the circumstances, his recalling that he intended to *C* would quickly result in his having an occurrent intention to *C*, the stronger his disposition to recall this, the more plausible the claim that he has a standing intention to *C*.

I will not attempt to construct an analysis of *standing intention*. However, after laying further groundwork, I will propose a sufficient condition for an agent's having a standing intention to *A* that applies to short-term standing intentions. If there are short-term standing intentions, representative examples include the following: Flori's standing intention to stop by the beer store on her way home from work today, Guy's standing intention to arrange to have flowers delivered to his partner today, and my standing intention to take Angela to a 7:00 movie tonight. In each case, the agent's forming or acquiring an occurrent intention to *A* plays a significant role in accounting for the existence of the disposition to have occurrent intentions to *A* in which the standing intention to *A* may consist. If I had never decided to take Angela to a 7:00 movie, nor otherwise acquired an occurrent intention to take her, I would not have a disposition to have occurrent intentions to do that; and similar things are true of Flori and Guy.

Might it be true that an agent's disposition to have occurrent intentions to *A* is a standing intention to *A* only if his having acquired an occurrent intention to *A* helps to account for the disposition's existence—a proposition that parallels a thesis about standing beliefs that I have discussed? Recall the case of Erwin and the hypnotist. If Al's disposition to acquire occurrent intentions to take Angela to a 7:00 movie—again, occurrent intentions to *M*, for short—is plausibly regarded as a standing intention to *M* whereas Erwin's disposition to acquire occurrent intentions to open a window whenever he hears the word “open” is not plausibly regarded as a standing intention to open a window, why is that? One difference between Al's disposition and Erwin's, given my suggestion about the former, is that Al's includes a disposition to recall that he intended to *M* whereas Erwin's includes no disposition to recall that he intended to open a window. A related difference is that Al's acquisition of an occurrent intention to *M* helps to explain why he is disposed to acquire occurrent intentions to *M* whereas Erwin acquired no comparable occurrent intention.

Suppose that Al's disposition to have occurrent intentions to *M* is realized in a neural state *N1*.²⁴ Suppose also that a supernatural being puts Al's friend Halle into a duplicate neural state *N2* without inducing in her an occurrent intention to *M*. She comes to be in *N2* before she has any occurrent intention to *M*. Does Halle, at this time, have a standing intention to *M*?

I think not. Al happens to have a genuine U.S. dollar bill in his pocket. Suppose that a supernatural being makes a duplicate of that bill and puts it in Halle's pocket. Did he put a genuine U.S. dollar bill in her pocket? Of course not: the bill is a counterfeit. Only bills suitably produced by the U.S. Treasury Department are genuine U.S. dollar bills. Suppose that Al has a sunburn on his right forearm and, as he sleeps, Halle produces a very similar burn on his left forearm with a heat lamp. Has she given Al a sunburn? Definitely not. Necessarily, a burn is a sunburn only if it is produced by exposure to the sun. My suggestion is that standing intentions, like genuine U.S. dollar bills and sunburns, have essential historical properties. If that is so, then, in one agent, a particular neural state may realize a standing intention to *A* whereas, in another, a duplicate neural state realizes a disposition to have occurrent intentions to *A* that is not a standing intention to *A*.

With this background in place, I float the following proposal:

SI. *S* has, throughout a span of time *t*, a standing intention to *A*, if (1) throughout *t*, *S* has a nondeviantly produced standing true belief that he acquired (or had) an intention to *A* and (2) he has nondeviantly come to be so constituted that the activation of this belief at any time during *t* will quickly generate an occurrent intention to *A*.

In some instances, the agent's having the occurrent intention that is generated is a matter only of its being a conscious non-standing intention to *A* (which

entails that at that time he intends to *A*). In others, it is at least partly a matter of the intention's being suitably at work in producing relevant intentional actions or items appropriate for the production of such actions. At 3:00, Al may be aware that he intends to take Angela to a 7:00 movie without the occurrent intention's doing any work of the kinds just mentioned at the time. If things work out, an occurrent intention to take Angela to a 7:00 movie will eventually help to produce appropriate actions.

Three technical points about the proposal are in order. The first concerns the expression "standing true belief." If a standing belief is wholly constituted by a disposition (of a certain kind) to have occurrent beliefs that *p*, it cannot literally be true; for dispositions cannot literally be true or false. By "S has a standing true belief that *p*," I mean the following: "S has a standing belief that *p*, and *p* is true."²⁵ Second, if (i) something is a standing belief that *p* only if it is a disposition to have occurrent beliefs that *p* such that the person's acquiring an occurrent belief that *p* played a role in producing the disposition, then (ii) *SI* is consistent with the truth of the claim that an agent's disposition to have occurrent intentions to *A* is a standing intention to *A* only if his having acquired an occurrent intention to *A* helps to account for the disposition's existence. If *i* is true and *S* has a nondeviantly produced standing true belief that he acquired an intention to *A*, *S*'s acquiring an occurrent intention to *A* is a cause of his acquiring an occurrent belief that he acquired an intention to *A* and his acquiring that belief played a role in producing the standing belief—in which case, his acquiring an occurrent intention to *A* played a role in producing the standing belief, thereby helping to account for the existence of his disposition to have occurrent intentions to *A* if the standing belief is part of that disposition. Third, an agent's having an occurrent distal intention to *A* at some time *ti* during *t* is compatible with his having at *ti* a disposition to acquire occurrent intentions to *A*, though it is not compatible with his having at *ti* a disposition to acquire at *ti* an occurrent intention to *A* (unless it is possible for an agent simultaneously to have two occurrent intentions to *A*: see n. 23). One cannot acquire at a time something that one already has at that time.

Another technical issue requires attention. Causal accounts of anything raise issues about deviant causal chains, and *SI* invokes the absence of causal deviance in two connections. Now, a deviant causal connection between an *x* and a *y* is deviant relative to *normal* causal routes from *x*-s to *y*-s. And when causal deviance is the issue, what counts as normal is perspective-relative. Here is a simple illustration. Halle, who is skilled in the use of firearms, intends to hit a certain target. She takes careful aim at it and fires. Her aim, surprisingly, is seriously errant. More surprisingly, her bullet ricochets off of several scattered objects into the target. From the point of view of physics, there is nothing strange or deviant about the causal sequence. But from the perspective of someone concerned to make judgments about intentional action, there is. Owing to the wayward causal chain linking Halle's intentional

firing to the bullet's striking the target, we say that hitting the target is not something she did intentionally. Readers are right to wonder how nondeviance should be spelled out in the context of *SI*. A full-blown attempt to spell this out would require, for example, a thorough examination of normal ways in which acquisitions of occurrent beliefs that *p* are involved in the production of dispositions to have occurrent beliefs that *p*. This issue, which has a major empirical component, is beyond the scope of this article. However, I have more to say about deviance shortly.

The activation of an agent's standing true belief that he acquired (or had) an intention to *A* might not have the result specified in *SI*, of course. An agent whose standing true belief that he acquired an intention to mop his kitchen floor tonight is activated when he arrives home from work may not acquire an occurrent intention to do that and may decide to postpone the chore. Mopping his floor tonight may now strike him as a bad idea, for example—or he may still regard it as a good idea but be unwilling to put it into action. (The kitchen is large and the agent is tired.) In another kind of case, an agent whose standing true belief that he acquired an intention to *A* tonight is activated may also recall that he subsequently decided against *A*-ing; he may not acquire an occurrent intention to *A*.

The former case points the way to a scenario of a kind that helps to account for the second mention of nondeviance in *SI*.²⁶ During breakfast, Ann decides to mop her kitchen floor tonight. Her day turns out to be exhausting—something she did not expect. In the past, when Ann has had an unexpectedly exhausting day, she has—quite reasonably—revoked her distal intentions to do tiring chores around the house after work when doing them was not urgent. But today, shortly before she leaves work, she is hypnotized, with the result that, if her standing belief that she intended to mop her floor tonight is activated, that event will immediately generate both an extremely powerful desire to mop her floor tonight and a false belief that it is very important to do that. Together, the desire and belief would remove any psychological obstacle to Ann's acquiring an occurrent intention to mop her floor tonight: such an intention would be quickly generated by the activation of the standing belief. If she had not been hypnotized, the activation of that belief would have understandably prompted Ann to decide to postpone the chore. If the claim that, in this scenario, after she is hypnotized, Ann has a standing intention to mop her floor tonight is counterintuitive, is *SI* threatened? No. Given the deviant way in which Ann comes to be so constituted that the activation of her standing belief will quickly generate an occurrent intention to mop her floor, this case is not a counterexample to *SI*.

Consider the following case. During breakfast, Al decides to mop his kitchen floor this evening. At 2:00, Bob calls to invite Al to a dinner party at a popular restaurant on the coast about fifty miles from Al's house. Al is not particularly fond of Bob and his friends. Nor is he fond of long drives,

especially after a long day's work. But Al is distracted when he receives the call, and, because he has skipped lunch, he is very hungry. His hunger makes Bob's proposal seem much more attractive than it otherwise would seem, and he—very uncharacteristically—accepts the invitation. Because Al knows that he will not have time to mop his floor before dinner and will not feel like doing it after the long drive home from the restaurant, he decides not to mop his floor tonight. Then he plunges back into his busy day at work. Soon, Al has a substantial snack at his desk. Not long after he finishes it, he comes to be psychologically so constituted that if, at any point during the remainder of his work day, his standing belief that he had an intention to mop his floor tonight is activated, that event will quickly generate two things: an occurrent intention to call Bob with an excuse for not attending the party and an occurrent intention to mop his floor tonight.

If *SI* is true, does Al have by 3:00 or so a standing intention to mop his floor tonight? That result is counterintuitive, given that Al has revoked his intention to mop his floor tonight; and it is easy to avoid. Augmenting *SI* with a third condition—namely, that *S* has not revoked his intention to *A*—would solve the problem. Similar cases can be constructed in which, although the agent does not *revoke* his intention to *A* (if that entails subsequently forming or otherwise acquiring an intention not to *A*), he does genuinely reopen the question whether to *A*. When an agent genuinely reopens this question, it is no longer true that he intends to *A*. The problem posed by both kinds of case can be blocked by adding the following condition to *SI*: *S* has not revoked his intention to *A* and has not genuinely reopened the question whether to *A*. I dub the augmented proposal *SI*^{*}.

Earlier, I remarked that a disposition to have occurrent intentions to *A* may be too weak to count as a standing intention to *A*. *SI*^{*} addresses the issue of strength in two ways. A disposition to have occurrent intentions to *A* that encompasses or depends on the standing belief featured in *SI*^{*} must meet whatever dispositional strength conditions obtain for standing beliefs, and condition 2 links activation of the standing belief tightly to the generation of an occurrent intention to *A*.

SI^{*} does not represent the activation of the agent's standing belief that he acquired an intention to *A* as doing all the work in generating an occurrent intention to *A*. The agent's more general psychological constitution matters. Must that constitution include a 3-intention to *A*? No. Even in the absence of a 3-intention, an agent who acquired an intention to *A* may be psychologically so constituted that, provided that he lacks stronger motivation not to *A* than to *A* upon acquiring the occurrent belief in which the activation of his standing belief consists, he quickly acquires an occurrent intention to *A*. A desire not to mop one's kitchen floor tonight—or to do something else instead of mopping it—may prevent one's recalling that one intended to do that from issuing in an occurrent intention to do it. In the absence of a significantly

strong desire of that kind—or, as one might say, in the presence of an *A*-friendly motivational condition—the agent may be so constituted that the recollection smoothly prompts the acquisition of an occurrent intention to mop the floor tonight.²⁷

In this article, I have developed a position on what it is for an intention to be an occurrent intention at a time and an account of the persistence of occurrent intentions to *A* as occurrent intentions to *A*, and I have defended a sufficient condition for something's being a standing intention to *A*. Although I focused on relatively short-term standing intentions, I have acquired an intention to explore their long-term counterparts elsewhere. If I have a standing intention to do that, I hope it will not need to persist very long.²⁸

Notes

¹ I occasionally use "*S* intends to *A*" as shorthand for "*S* has an intention to *A*." Even some readers who believe that there are intentions may regard this as a potentially misleading practice. They may think that whereas (1) "Al intends to finish sweeping his floor and then put the corkscrew away" entails (2) "Al intends to put the corkscrew away," (3) "Al has an intention to finish sweeping his floor and then put the corkscrew away" does not entail (4) "Al has an intention to put the corkscrew away." The thought is that the latter entailment would call for two states of mind—an intention to finish sweeping the floor and then put the corkscrew away and a separate intention to put the corkscrew away—whereas the former does not require more than one. However, a philosopher who claims that 3 entails 4 might mean nothing more than that, necessarily, the state of mind reported in 3 includes the one reported in 4.

² See Brand 1984, Enç 2003, and Mele 1992.

³ See Mele 1992, chs. 7–13 for an account of the nature and functions of nondispositional intentions.

⁴ For the distinction, see, for example, Brand 1984, Bratman 1987, and Mele 1992, 2003. The terminology varies.

⁵ In the remainder of this paragraph and in the next one, I borrow from Mele 2003, 30–32.

⁶ For discussion of some relevant empirical work, see Vollmer 1993.

⁷ Gilbert Harman writes, "one believes something explicitly if one's belief in that thing involves an explicit mental representation whose content is the content of that belief"; later, he speaks of something "explicitly represented in one's mind" as "written down in Mentalese as it were, without necessarily being available to consciousness" (1986, 13–14). This is the general picture I mean to evoke with the expression "explicit attitudes," but I neither venture an opinion about *how* the contents of attitudes are explicitly represented nor claim that they are all represented in the same way (e.g., in a language of thought). Incidentally, my subsequent use of "implicit" in "implicit attitudes" does not parallel Harman's use of the term in his expression "implicit beliefs" (13).

⁸ On the question whether standing desires are actually desires, see Mele 2003, 32–33.

⁹ In an atypical case, an agent dies a millisecond after he acquires a proximal intention to raise his right arm. If agents do not begin to execute intentions for arm raisings for some milliseconds after the intention is acquired, this agent did not begin to execute his intention. Incidentally, I am not assuming here that agents do not begin executing their intentions for overt actions until muscles begin to move. On this conceptual issue, see Adams and Mele 1992.

¹⁰ See Mele 2003, ch. 9. Also see Frankfurt 1988, 174–76, McCann 1986, 254–55, Pink 1996, 3, and Searle 2001, 94. I am not suggesting that all intentions are formed in acts of deciding. Indeed, as I have observed elsewhere, many of our intentions seem to be acquired without being

formed in acts of deciding. For example, “When I intentionally unlocked my office door this morning, I intended to unlock it. But since I am in the habit of unlocking my door in the morning and conditions... were normal, nothing called for a *decision* to unlock it” (Mele 1992, 231). If I had heard a fight in my office, I might have paused to consider whether to unlock the door or walk away, and I might have decided to unlock it. But given the routine nature of my conduct, there is no need to posit an action of intention formation in this case. My intention to unlock the door may have been acquired without having been actively formed. If, as I believe, all decisions about what to do are prompted partly by uncertainty about what to do (Mele 2003, ch. 9), in situations in which there is no such uncertainty, no decisions will be made. This is not to say that, in such situations, no intentions will be acquired.

¹¹ For other grounds for believing that occurrent intentions may produce actions at times at which the agent is not conscious of the intentions, see Marcel 2003 and Mele 2004.

¹² Suppose that an agent has a standing intention to buckle his seat belt before starting his car. In playing the role characteristic of standing intentions, a role that requires its being manifested in an occurrent intention to buckle his seat belt before starting his car, this intention may be said to play an indirect role in initiating his buckling his seat belt now.

¹³ As I understand awareness, it is an essentially occurrent phenomenon.

¹⁴ I use the word “suitably” to signal two things: first, that the operative sense of “at work” is one that requires that the work at issue is not the characteristic work of standing intentions; second, the exclusion of pertinent kinds of monkey business (for example, a being with special powers may detect a nonoccurrent distal intention to *A* that I have and respond to the persistence of that intention by causing me, as it persists as a nonoccurrent intention, to begin thinking about how to put myself in a position to *A*). For my purposes in this article, a more rigorous statement of condition 1 is not required.

¹⁵ In some circumstances, an intention’s producing a desire for something *M* that the agent takes to be a means to executing that intention may prevent him from executing the intention. Even so, on the broad reading of “items appropriate for the production of relevant intentional actions” at work in 1 here, the desire for *M* is such an item. The appropriateness, in this case, is subjective.

¹⁶ Anthony Marcel (2003) criticizes the view of some psychologists (e.g., Wegner 2002) that intentions are essentially conscious. For a discussion of various interpretations of “conscious intention,” see Mele n.d., where a more refined account of conscious intentions than the one suggested by condition 2 is developed. For the purposes of the present article, motivating the refinements would be an unnecessary—and lengthy—distraction.

¹⁷ In this paragraph and the following three, I borrow from Mele 2003, 27–28.

¹⁸ The guidance depends on the agent’s monitoring progress toward his goal. The information (or misinformation) that Jan has entered her password, for example, figures in the etiology of her continued execution of her plan. On guidance, see Mele 2003, 55–62.

¹⁹ In the case of an intention for a not-doing (e.g., an intention not to vote tomorrow), the agent may instead be settled on not violating the simple plan embedded in it—the plan not to vote. On not-doings and attitudes toward them, see Mele 2003, 146–54.

²⁰ A critic may claim that in all cases of this kind the agent is settled on a course of action without realizing it and that he is deliberating only because he does not realize what he is settled on doing. For argumentation to the contrary, see Mele 1992, ch. 9.

²¹ Charles Hermes and Michael Robinson persuaded me (in conversation) of the plausibility of this view.

²² If, from 11:59 until noon, one is conscious of one’s intention to *A* when the noontime bell begins to toll, is one conscious of it at each nanosecond during that minute? If the nature of consciousness is such that the answer is *no*, then “time” in *OIPb* should be understood as something the units of which are appropriate for the timing of conscious experiences. The same goes for the timing of an intention’s being “at work.” (Discussion with Charles Hermes motivated this note.)

²³ If it is conceptually possible for an agent to have, at the same time, two different occurrent intentions with the same plan component, there is a related problem with *OIPb*. Condition 2 may be satisfied even if *NI* ceases to exist before *t2*. I avoid a potentially lengthy technical discussion by stipulating that if intention pairs of this kind are conceptually possible, *OIPb* is to be understood as restricted to agents who have no such pairs of intentions. (Readers who believe that agents often have intentions to *A* while also having distinct intentions to try to *A* and that these intentions have the same plan component will see the restriction as severe. Such readers may find this view attractive partly because they believe that intentionally *A*-ing (often) involves trying to *A* and that the *A*-ing and the trying to *A* are produced by distinct intentions. My own view is that in standard cases of agents who both intend to *A* and try to *A*, the intention to *A* is at work in producing the trying in the absence of any distinct intention to try to *A*.)

²⁴ Here I have in mind what I elsewhere termed a “narrow” neural state: “*S*’s narrow [physical] condition at *t* . . . is a condition that *S* would share with any molecule-for-molecule twin. It is an abstraction consisting of *S*’s broad condition stripped of its historical properties and any other relational properties involving *S*-external relata” (Mele 1992, 23–24).

²⁵ In the case of occurrent beliefs that *p*, if occurrent beliefs are states of mind, to say that I have an occurrent true belief that *p* is not to say that the state of mind is true. Rather, it is to say that I have an occurrent belief that *p* and that what I believe—namely, *p*—is true.

²⁶ The first appeal to nondeviance is intended to ensure that the agent’s standing true belief that he acquired an intention to *A* has his acquiring an occurrent intention to *A* as a cause.

²⁷ Whether 3-intentions exist is an empirical question, as I have said. If they do, they are, in principle, linked to standing and occurrent intentions in interesting ways (if intentions of the latter two kinds also exist). The presence of a 3-intention to *A* in an agent can be a partial ground of his having a standing intention to *A*, as I have mentioned. And a 3-intention to *A* can become an occurrent intention to *A* and can persist while it has the property of being an occurrent intention to *A*. In principle, the same 3-intention may acquire that property on more than one occasion. It may become occurrent, cease being occurrent, and become occurrent again.

²⁸ For comments on a draft of this article, I am grateful to Randy Clarke, Charles Hermes, Michael Robinson and two anonymous referees.

References

- Adams, Frederick, and A. Mele. (1992) “The Intention/Volition Debate” *Canadian Journal of Philosophy* 22: 323–38.
- Alston, William. (1967) “Motives and Motivation,” in P. Edwards, ed., *The Encyclopedia of Philosophy*, New York: Macmillan.
- Audi, Robert. (1994) “Dispositional Beliefs and Dispositions to Believe” *Noûs* 28: 419–34.
- Brand, Myles. (1984) *Intending and Acting*, Cambridge, Mass.: MIT Press.
- Bratman, Michael. (1987) *Intention, Plans, and Practical Reason*, Cambridge, Mass.: Harvard University Press.
- Enç, Berent. (2003) *How We Act*, Oxford: Clarendon Press.
- Frankfurt, Harry. (1988) *The Importance of What We Care About*, Cambridge: Cambridge University Press.
- Goldman, Alvin. (1970) *A Theory of Human Action*, Englewood Cliffs: Prentice-Hall.
- Harman, Gilbert. (1986) *Change in View*. Cambridge, Mass.: MIT Press.
- Marcel, Anthony. (2003) “The Sense of Agency: Awareness and Ownership of Action,” in J. Roessler and N. Eilan, eds., *Agency and Self-Awareness*, Oxford: Clarendon Press.
- McCann, Hugh. (1986) “Intrinsic Intentionality” *Theory and Decision* 20: 247–73.
- Mele, Alfred. (1992) *Springs of Action*, New York: Oxford University Press.
- Mele, Alfred. (2003) *Motivation and Agency*, New York: Oxford University Press.

- Mele, Alfred. (2004) "The Illusion of Conscious Will and the Causation of Intentional Actions" *Philosophical Topics* 32: 193–213.
- Mele, Alfred. (n.d.) "Conscious Intentions," in J. Campbell, M. O'Rourke, and H. Silverstein, eds., *Action, Ethics, and Responsibility*, Cambridge, Mass: MIT Press (forthcoming).
- Pink, Thomas. (1996) *The Psychology of Freedom*, Cambridge: Cambridge University Press.
- Searle, John. (1983) *Intentionality*, Cambridge: Cambridge University Press.
- Searle, John. (2001) *Rationality in Action*, Cambridge, Mass: MIT Press.
- Vollmer, Fred. (1993) "Intentional Action and Unconscious Reasons" *Journal for the Theory of Social Behavior* 23: 315–26.
- Wegner, Daniel. (2002) *The Illusion of Conscious Will*, Cambridge, Mass.: MIT Press.